

Checkpoint 2 Rubric Review

Alignment check: `notebooks/index.py` and `preso/checkpoint2_deck.md` against the Checkpoint 2 Evaluation Rubric

Team .egsy intelligence (Group 14)

August 2, 2026

Summary

Overall alignment against the Checkpoint 2 Evaluation Rubric is strong. Seven of nine graded categories are fully met at the “Excellent” tier. Two items still need attention before submission: an asymmetry between the two model proposals’ tooling/compute detail, and a comparison-of-models section that is currently distributed rather than consolidated. The presentation’s speaker-notes typo has been fixed, and the deck’s timed run-through clocks in at 14 minutes, within the 15-minute budget.

Category-by-Category Findings

Problem Context (Steps 1-2 Summary) — 5 pts — Excellent

`checkpoint_2.py`’s opening cell explicitly bridges from Step 2 (“With our data cleaned and scored in Step 2, we turn here to...”), and via `index.py`’s `embed/vstack`, the full `checkpoint_1.py` (Steps 1-2) precedes it in the combined write-up — not just a summary, the whole thing.

Step 3: Evaluation Metrics — 15 pts — Excellent

Per-class precision/recall/F1, recall-as-constraint, macro-F1 as the ranking metric, and a 3x3 confusion matrix, each with an explicit rationale (`checkpoint_2.py`, “Why plain accuracy is the wrong headline metric” -> “Metric framework”). The deck’s “Metrics & Success Thresholds” slide mirrors this with matching numbers.

Step 3: Success Criteria — 10 pts — Excellent

RECALL_FLOOR / MACRO_F1_FLOOR / PRECISION_FLOOR (0.70 / 0.60 / 0.45), each benchmarked against a random-uniform baseline and applied identically to both models.

Step 3: Evaluation Methodology — 15 pts — Excellent

Study-grouped split, exhaustive search vs. `StratifiedGroupKFold` benchmark, and McMahon held out with stated reasoning — `checkpoint_2.py`’s “Split strategy - group by study” section and the deck’s “Split Strategy” slide agree on the numbers (190/46 sites, 19.5% / 3.5%). A recent pass trimmed the split-comparison preview table from 20 to 5 rows and combined count/percent columns for print-width control; the deck never cited the dropped rows, so no narrative inconsistency resulted.

Step 4: Model Proposal #1 — 15 pts — Excellent

Model A has a dedicated “Tooling & compute plan for baseline” section — specific libraries, an explicit operation-count estimate, hardware needs, and reproducibility — plus the technique / rationale / training-plan / strengths-limitations / suitability structure.

Step 4: Model Proposal #2 — 15 pts — Good, not yet Excellent

Model B (“Model B: random forest classifier”) has the same technique/rationale/training-plan/strengths-limitations/suitability structure, but **still has no equivalent standalone tools/compute- estimate section**. Its “Training and optimization plan” mentions `n_jobs=-1` and “standard CPU, no GPU” in passing, but nothing like Model A’s operation-count math or explicit software list. The rubric explicitly asks for “tools, computational needs” on each proposal — this remains the one place the two are not symmetric.

Suggested fix: add a short “Tooling & compute plan” subsection for Model B mirroring Model A’s, with a rough operation-count or wall-clock estimate for the random forest grid search.

Comparison of Models — 10 pts — Good, not yet Excellent

Trade-offs are covered, but only in a distributed way — each proposal’s own “Expected strengths and limitations” cross-references the other, and the deck’s “Two Competing Proposals” / “Validation & Trade-offs” slides add a top-level framing. There is still no single dedicated section in `checkpoint_2.py` that pulls scalability, complexity, expected performance, and a stated preference into one place. Not naming a preferred model is correct at this stage (no Step 5 results exist yet), but a consolidated comparison section would better match the rubric’s “Excellent” tier.

Suggested fix: add a short “Comparison of Models” section to `checkpoint_2.py` (before or after the two proposals) that explicitly restates the scalability/deployment metric, complexity trade-off, and expected-performance framing already present elsewhere, in one place.

Individual Contributions — 5 pts — Excellent

Both `index.py`’s roles table/prose and the deck’s “Team & Roles” slide are specific and consistent with each other.

Organization & Technical Writing — 5 pts — Excellent

The speaker-notes typo in `preso/checkpoint2_deck.md` (“...contain no personally identifiable information. so there are no privacy concerns...”) has been fixed to a single, correctly-punctuated sentence. A recent table-trimming pass across `checkpoint_1.py`, `checkpoint_2.py`, and `footer.py` removed decorative/duplicative tables (sample-identifier previews, redundant missing-value breakdowns) and combined count/percent columns, keeping wide tables within print width without losing any of the underlying suitability assessments. Heading levels are consistent with this project’s current structure: each checkpoint notebook and `footer.py`’s Conclusion use the same top-level header weight, matching how the document is composed across multiple files rather than a single notebook.

Presentation Quality (15 min) — 5 pts — Excellent

The deck's timed run-through clocks in at 14 minutes, within the 15-minute budget, with every team member presenting a distinct section.

Outstanding Action Items

1. Add a “Tooling & compute plan” subsection for Model B in `checkpoint_2.py`.
2. Add a consolidated “Comparison of Models” section in `checkpoint_2.py`.